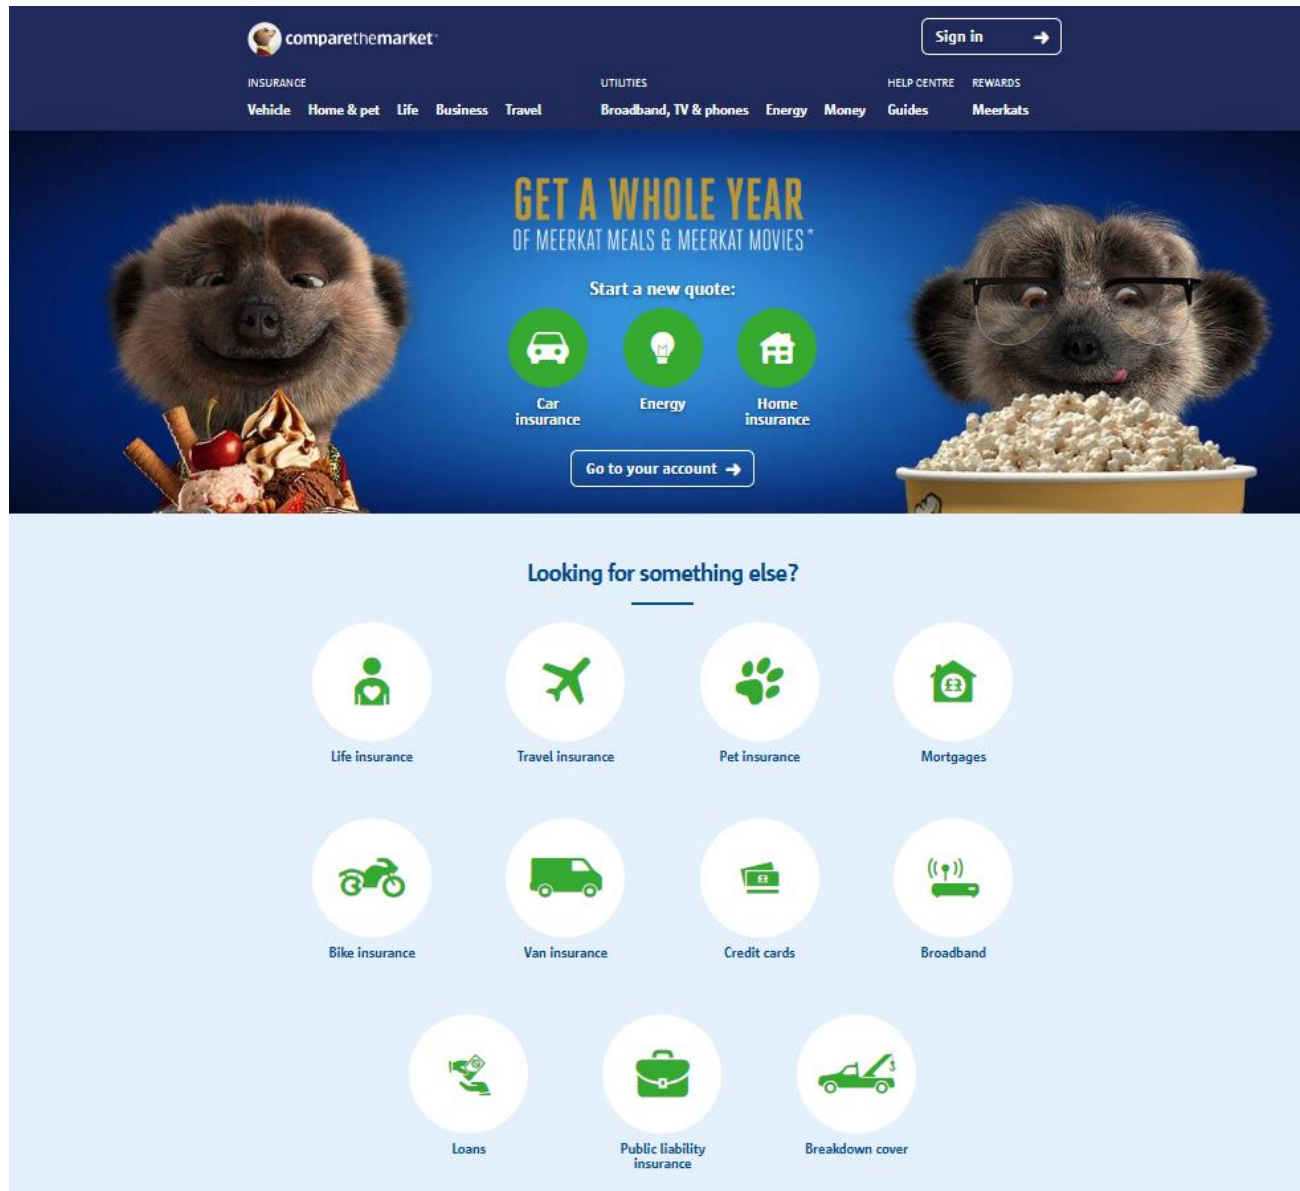


- Why microservices?
- Essence of microservices
- Refactoring microservice-based architectures
- Concluding remarks
- Case studies
  - [comparethemarket.com](http://comparethemarket.com)



UK's most popular (insurance) price comparison website



**comparethemarket.com**



# comparethemarket.com



- *Price comparison web site, various product verticals*
- *Initially one team, one system*
- *Rapid growth of customer numbers and staff*
- *Change*
  - *from monolith architecture to microservices with single responsibility*
  - *from relational DB (point of contention, changes breaking other teams'apps) to Docker-based approach with mongoDB EE*
- *Achievement*
  - *ability to change quickly to catch new consumer trends and tech adoption trends*

# comparethemarket.com

- Issues of monolith
  - coordinate releases amongst teams
  - long feature cycles
  - changes negatively impacting other product teams
  - not being able to expose their functionality for partners through APIs
  - inability to scale a single database with 200+ tables and a single point of failure any other way but vertically

- Solution

Change the layout of the internal teams from groups of specialists (UI, DBA, MW) to **autonomous product teams** that could change what was theirs through **microservices**, a set of smaller applications chained together instead of a single, master application

# Benefits

## 1 More manageable

“By taking a single complicated thing and breaking it into smaller things, they become easier to reason about.” (Unix design philosophy)

## 2. Flexible tech stack

“We haven’t had to commit to a tech stack, we use the right tools for the problem e.g. if we have a machine learning problem we use Python as there is a wealth of open source libraries we can make use of. If we just want to do a restful API we use Node.js.”

## 3. Reduced cost of failure

“Big services tend fail in a big way. If your monolithic application isn’t working you aren’t making money. In a microservice world if a service stops working it doesn’t take down the entire businesses ability to make money. It drives innovation and shortens the build, measure, learn cycle so people feel more empowered to experiment and it is freeing to build things that don’t have to last ten years.”

## 4. More specialized

“You can use microservices to specialise what you do. So if you have customer data you can put that into a single microservice, wrap that in layers of additional security and not burden the rest of your services with that same constraint. In a monolithic world there needs to be that highest common denominator.”

## 5. Independence

“Teams can have their own backlog of change and scale and release independently which allows the organisation as a whole to move faster.”

# Drawbacks

## 1. Complexity

“By splitting one large thing into small bits you end up having more bits.”

## 2. Disparate data sources

“You end up with many data sources and so how do you deal with that data?”

## 3. Creating a distributed monolith

“You have to put some effort into design upfront. Just taking a monolith and breaking it into microservices won’t cut it if those bits were tightly coupled to start with, you will just end up with tightly coupled microservices, which is a distributed monolith, which is the worst of all worlds.”

To minimise these drawbacks:

- **investment in automation**

”Trying to realise **continuous delivery** by making releases reliable and repeatable and the whole process a non-event, rather than a 4am in the morning cold sweat process.”

- **ensure that all monitoring and alerting frameworks looked the same**

“all of our microservices emit a standard set of metrics for latency, throughput etc. All of these metrics have a standardised threshold.”

- **realistic approach to faults and failures**

“by accepting the fact that computers break and networks fail we have focused on becoming fault tolerant.”

if a non-critical dependency goes down a service will still **respond in a degraded form**

if the failure is a critical dependency like the database, it **fails fast** “so that the client doesn’t expend lots of resources waiting for a response that doesn’t come back”

- **data management**

“Each microservice emits data in the form of a JSON object. This gives us a real-time view of what is happening across our estate and is scalable.”

“Each microservice has a private data store so no one else relies on the structure of the data which aligns with the Mongo no-schema metadata approach. The **team can change the structure of the data in that store without having to coordinate.**”



- Why microservices?
- Essence of microservices
- Refactoring microservice-based architectures
- Concluding remarks
- Case studies
  - [comparethemarket.com](http://comparethemarket.com)
  - Spotify



PLAYLIST

# Driving

Pop jams for the car

Created by: **Ari Vaniderstine** • 28 songs, 1 hr 38 min

**PLAY** ⋮

FOLLOWERS 0

Go to Playlist Radio

Collaborative Playlist

Make Secret

Edit Details

Report

Delete

Create Similar Playlist

Download

Share

TITLE	ALBUM		Download	
+ Shut Up and Dance	TALKING IS HARD	2015-08-09	3:19	
+ Cough Syrup	Young the Giant (Special Edit...	2015-08-09	4:10	
+ Pumped Up Kicks	Torches	2015-08-09	4:00	
+ Take a Walk	Gossamer	2015-08-09	4:24	
+ Work This Body	TALKING IS HARD	2015-08-09	2:56	
✓ Radioactive	Night Visions	2015-08-09	3:07	
+ Everybody Talks	Neon Trees	Picture Show (Deluxe Edition)	2015-08-09	2:57
+ Little Talks	Of Monsters and Men	My Head Is An Animal	2015-08-09	4:27
+ Little Lion Man	Mumford & Sons	Sigh No More	2015-08-09	4:05
+ Geronimo	Sheppard	Bombe Away	2015-08-09	3:38
+ I Will Wait	Mumford & Sons	Babel	2015-08-09	4:37

## Listen Offline

On Premium

1. Tap **Your Library**
2. Select from **Playlists, Songs, Albums, Artists, Podcasts & Videos** (for podcasts)
3. Select one
4. On a **WiFi** connection ...  
Switch **Download** on
5. Play everywhere  
Even without internet!

For more help visit [support.spotify.com](https://support.spotify.com)

# Some requirements

- Scale to **millions of users**
- Support **multiple platforms**
  - Web, mobiles, laptops, embedded devices, game platforms, ...
- Handle **complex business rules**
- **Be competitive** in a fast moving market
  - React quickly
  - Innovate

# Some numbers

- 75M+ Monthly Active Users
  - Long user sessions (average 23 mins), white noise all night
- 58 countries
- >20K songs added per day
- >2B user-generated playlists (+ 75M playlists created by Spotify)
- Incredibly complex business rules
  - choice of which version of track is served depends on business rules
  - different licensing agreements & contracts across different countries and labels
- Lots of competition



How do you support  
these requirements  
while moving fast  
and innovating?

# Solution

Autonomous full-stack teams


## Autonomous

adjective  
au-ton-o-mous - /ô-'tô-ne-mas/

(of a country or region) having the freedom to govern itself or control its own affairs.  
"the federation included sixteen autonomous republics"

**having the freedom to act independently.**

"school governors are legally autonomous"  
synonyms: self-governing, independent, sovereign, free, self-ruling, self-determining, autarchic;  
self-sufficient  
"an autonomous republic"



- Standard way of building connected apps
- «Horizontal» teams
- Hard to build new feature: lots of asking to other teams

*Client UX implementation*

*Core Library implementation*

*Server implementation*



*Infrastructure implementation*

*depends on*

*depends on*

*depends on*



## Full-stack autonomous teams




## Full-stack autonomous teams



Requires you to structure your application in loosely coupled parts






Datastore slow ...








Server is slow ...



Scaling?  
Changing?




# Microservices yay!

- **Easier to scale** based on real-world bottlenecks
- **Easier to test**
- **Easier to deploy**
- **Easier to monitor**
- Can be **versioned independently**
- Are **less susceptible to large failures**



*Crucial for embedded devices.  
Manufacturers will not update device  
ever, lamp will always use (old) API  
version*



*Users do not realize  
if «suggest service» fails*


# Microservices boo!

- Are harder to monitor

*Thousands of instances running together*

- Need good documentation / discovery tools

*Many services named by developers out of context*



- Create increased latency

*Lots of services calling other services*



*Netflix view aggregation model*

- Microservices for years now
- Production environment running in Spotify DCs (test environments run on AWS)
- Super happy of architecture
  - great for **scaling**
  - great for **rewriting** when needed
  - developers **deploying** by themselves whenever they want to



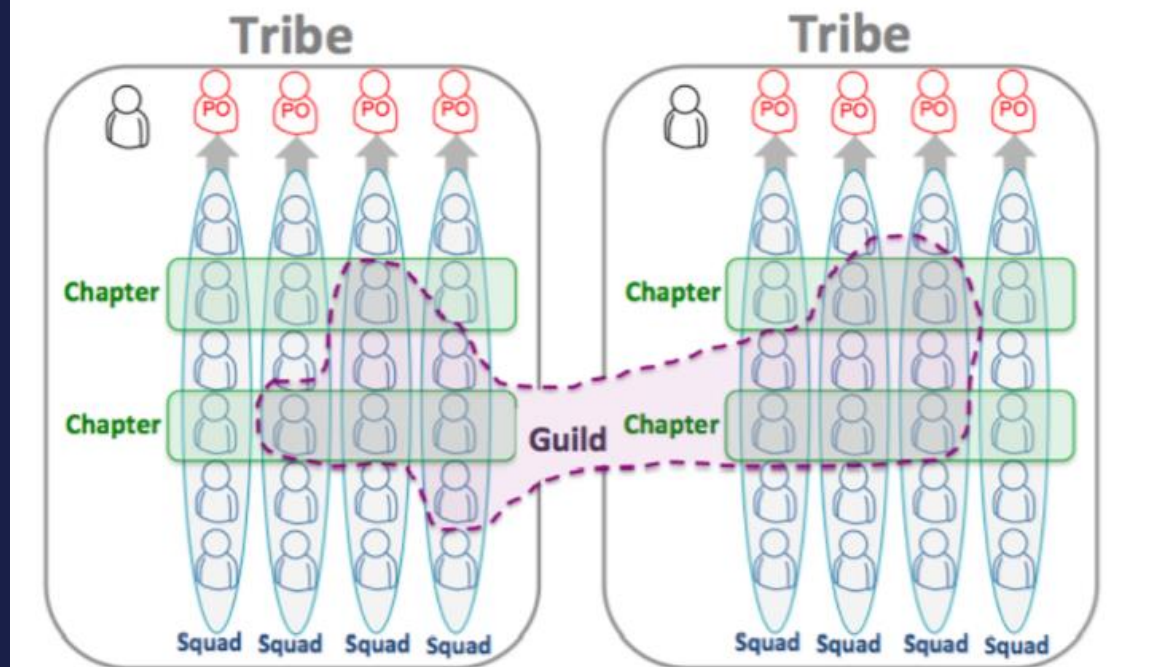
### Some numbers

- 90+ teams (squads)
- 600+ developers
- 5 development offices in 2 continents
- 1 product!

- 810 active services
- ~10 services per squad

Q: «squads»?

A: Agile team organization ...



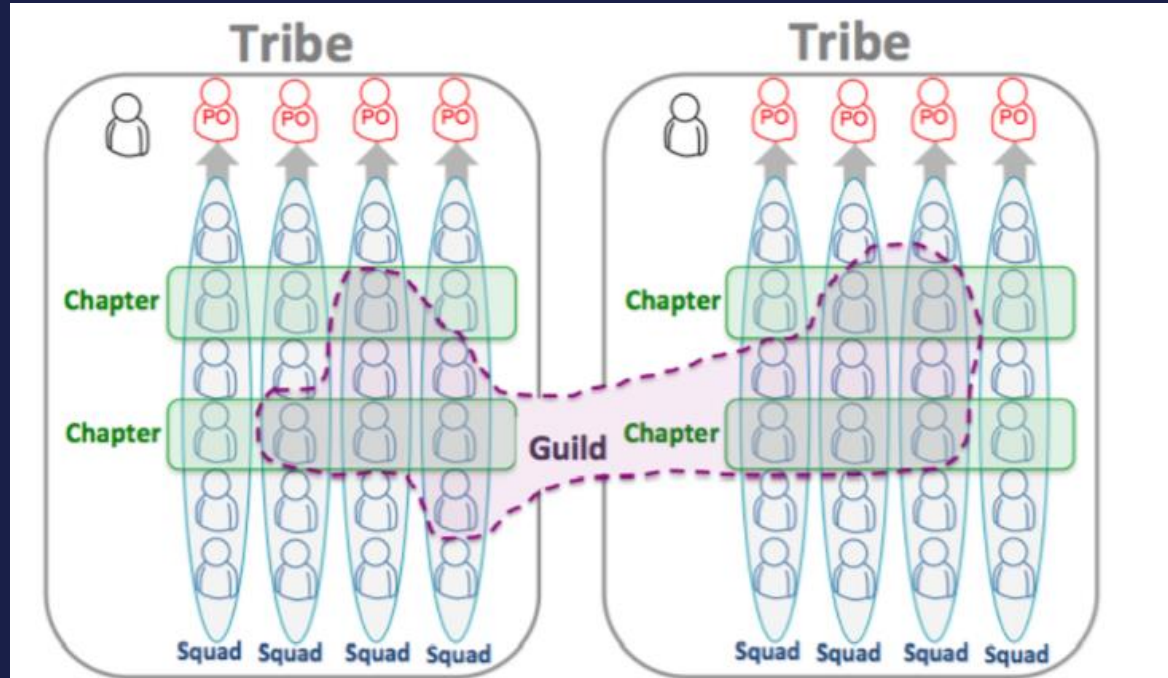
A **squad** has a dedicated Product Owner who feeds user stories to build. Squads have all the skills and tools needed to design, develop, test and release to production, being an autonomous, self-organising team, experts in their product area.

A **tribe** is a collection of squads within the same business area (e.g., mobile). The squads within a tribe sit in the same area. Spotify implements shared lounges to create inter-squad interaction, and regular informal team events and get together where squads share what they are working on. Tribe Leader is responsible for providing the right environment for all the squads.

**Chapters** are team members working within a special area (e.g., front office developers, back office developers, database admins, testers). Chapter members exchange ideas to promote innovation and 'cross pollination' across teams.

A **guild** is a community of members across the organization with shared interests (e.g., web technology, test automation), who want to share knowledge, tools code and practices.





Q: How is code reviewing performed at Spotify?

A: Chapters (sometimes guilds) do code reviews for squads.

Two «+1» required to merge.



# Spotify engineering culture

rent

Limited Blast Radius

via Decoupled Architecture



via Gradual Rollout



Experiment

A or B?

Let's try both and compare

If everything is under control, you're going too slow!

- Mario Andretti



Hack Time

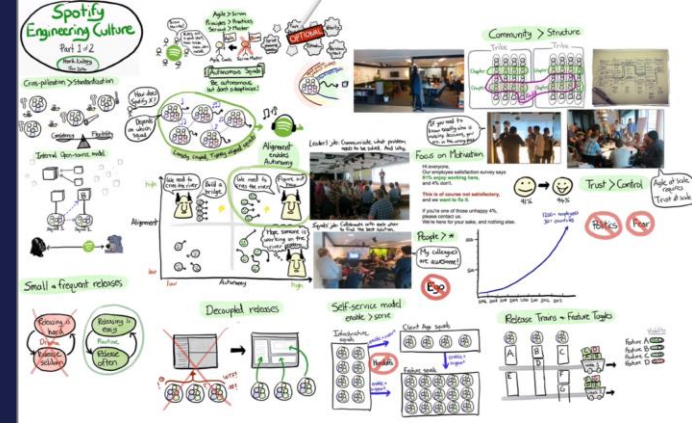
≈ 10%



People are not

## [Part I]

- Agile principles
  - Loosely coupled, tightly aligned squads
  - Cros- pollination
  - Internal open-source model for code
  - Frequent releases, enabled by decoupling
- Focus on people: motivation, community, trust



- Fail fast → learn fast → improve fast
  - Spotify is a fail-friendly environment
    - Fail walls, post-mortems, retrospectives
  - Failure recovery (rather than failure avoidance)
  - Limited blast radius via decoupled architecture and gradual rollout
- Product development approach
  - Lean startup principles (think | build | ship | tweak )
  - Impact > Velocity
  - Innovation > Predictability
  - Value delivery > Plan fulfillment
- 10% hack time
- Experiment-friendly culture

- Waste-repellent culture
  - «if it wroks, keep it, otherwise dump it»
  - Minimize need of big projects
  - Minimal bureaucracy to avoid chaos
  - Improvement boards, Toyota improvement «kata»
- «Health culture heals broken process»
  - Culture-focused roles (agile coaches)
  - Boot camps
  - Story telling

- Why microservices?
- Essence of microservices
- Refactoring microservice-based architectures
- Concluding remarks
- Case studies
  - [comparethemarket.com](http://comparethemarket.com)
  - Spotify
  - Netflix



## NETFLIX

Leader in subscription internet tv service  
Hollywood, indy, local  
Growing slate of original content

86 million members  
~190 countries, 10s of languages  
1000s of device types

Microservices on AWS

NETFLIX ORIGINALS >





# GLOBAL APPLICATION TRAFFIC SHARE

1

NETFLIX

14.97%



2.92%





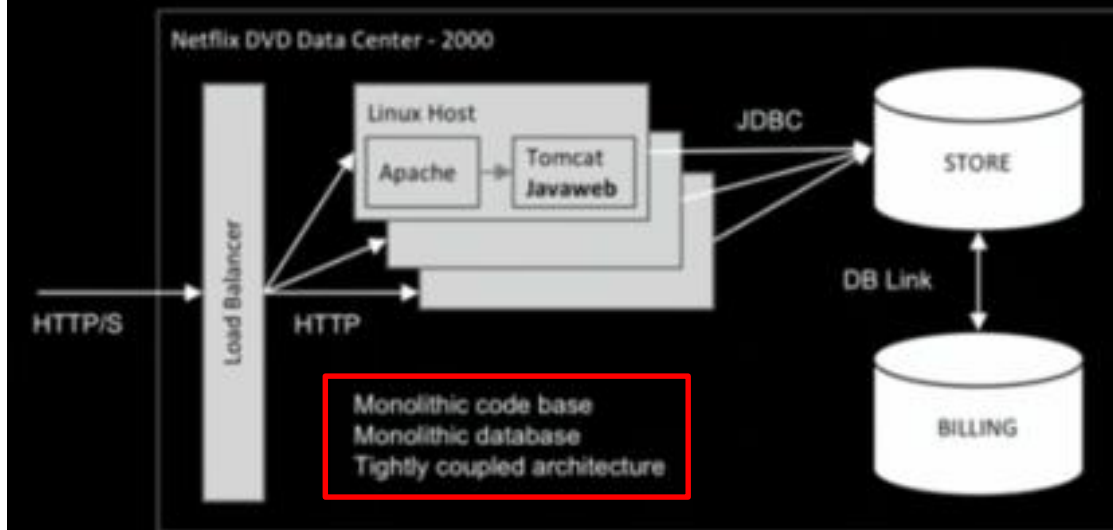


- Netflix figures
  - 86M members
  - 150M hours of streaming per day
  - 100,000+ AWS instances
- Migration to AWS
  - [2008, 2016]
  - improved productivity and scalability
- 500+ microservices
- Many tools
  - Spinnaker (CI/CD)
  - Atlas (monitoring)
  - Simian Army (testing reliability)
  - ...
- Shift to containers
  - Titus (container management environment)
  - then AWS Blox ...

- **Microservices**



## What microservices are not



Hard to scale  
Hard to modify

# What is a microservice?

*...the microservice architectural style is an approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API.*

*- Martin Fowler*

# What is a microservice?

## An Evolutionary Response

### Separation of concerns

Modularity, encapsulation

### Scalability

Horizontally scaling

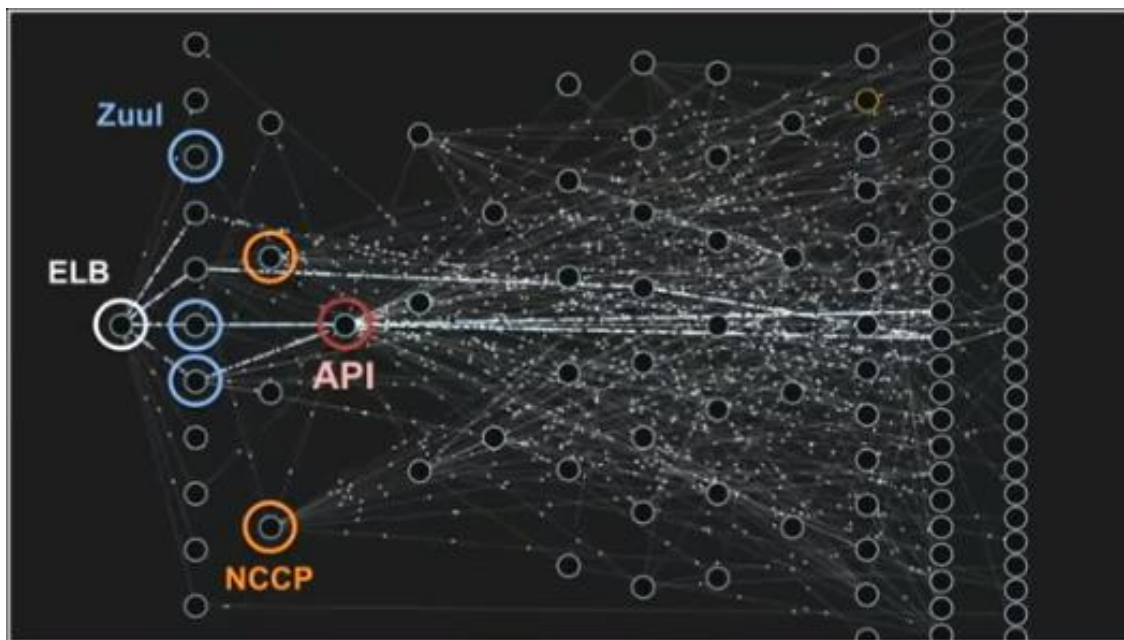
Workload partitioning

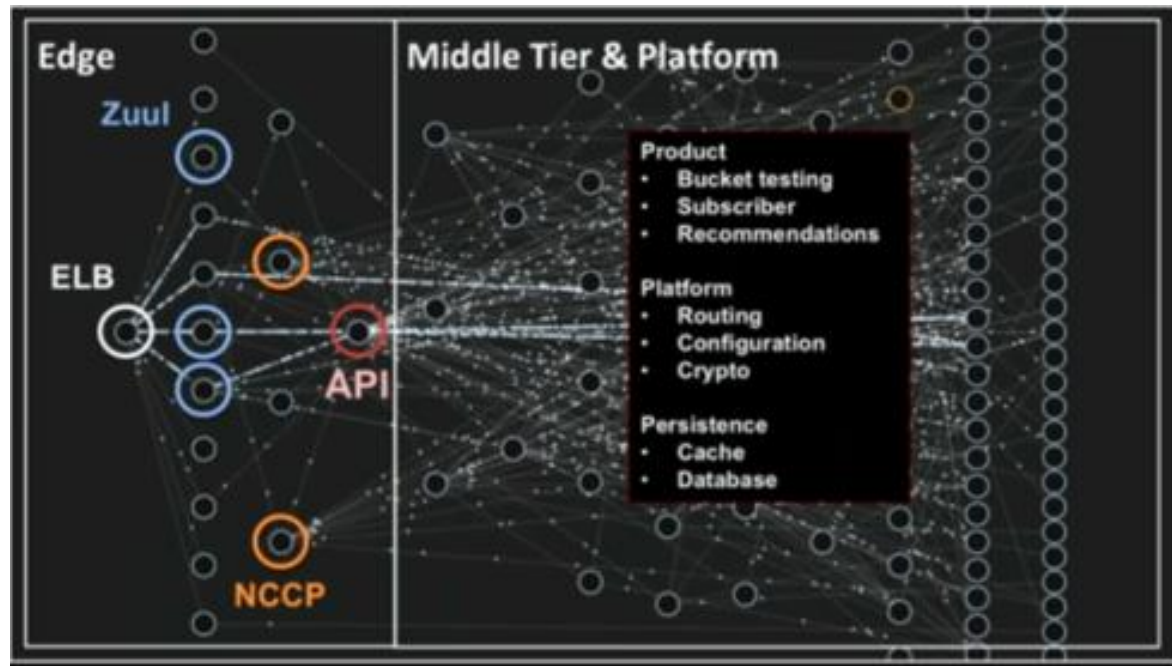
### Virtualization & elasticity

Automated operations

On demand provisioning







- **ELB**: AWS Elastic Load Balancing
- **Zuul**: proxy layer, performs dynamic routing
- **NCCP**: legacy tier, supporting earlier devices
- Netflix's **API** gateway, calling all other services

# Microservices are an abstraction

The diagram illustrates the abstraction of microservices. It shows a Client Application containing a Client Library, which is divided into an EVCache Client and a Service Client. The EVCache Client interacts with multiple EVCache instances (represented by cylinders with 'E' and 'V'). The Service Client interacts with multiple Service instances (represented by squares with 'S'), which in turn interact with multiple Database instances (represented by cylinders with 'DB').

... microservices are not always small

- Microservices
- Challenges
  - **Dependency**
  - Scale
  - Variance
  - Change

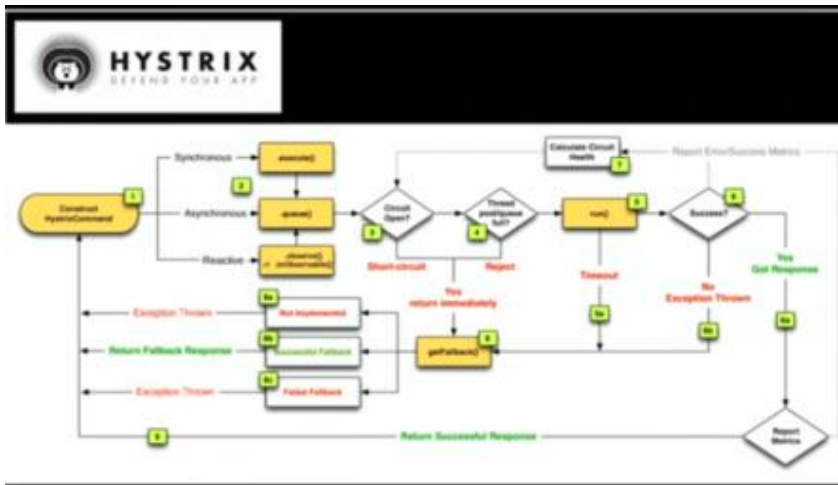
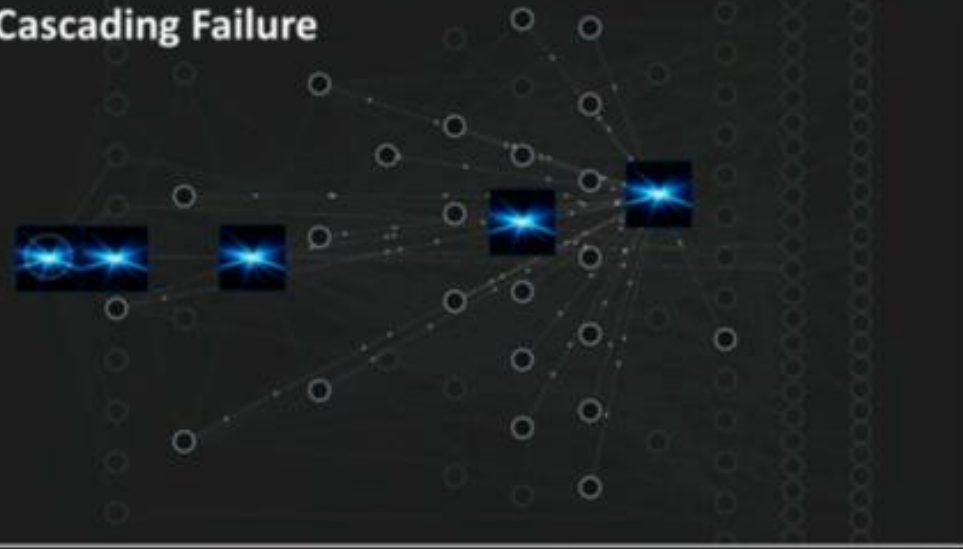
# Dependency: Inter-service requests

## Crossing the Chasm



Network latency, congestion, failure  
Logical or scaling failure

## Cascading Failure




Timeouts and retries  
Fallback (to allow customers to continue)  
Isolated thread pools and circuits



# Dependency: Inter-service requests (cont.)


How do you know if it works?

## Fault Injection Testing (FIT)



Synthetic transactions  
Override by device or account  
% of live traffic up to 100%

## Fault Injection Testing (FIT)

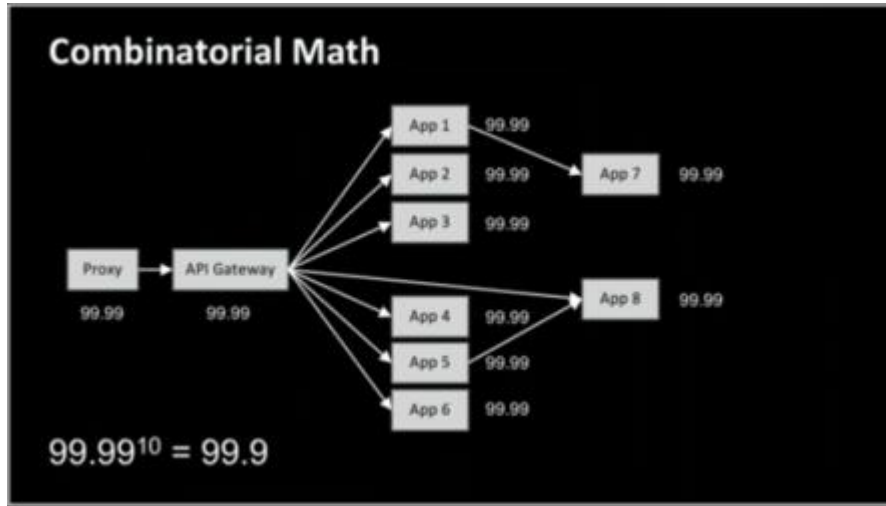


Enforced throughout the call path


**NETFLIX**  
**SIMIAN**  
**ARMY**





## How to constrain testing scope?



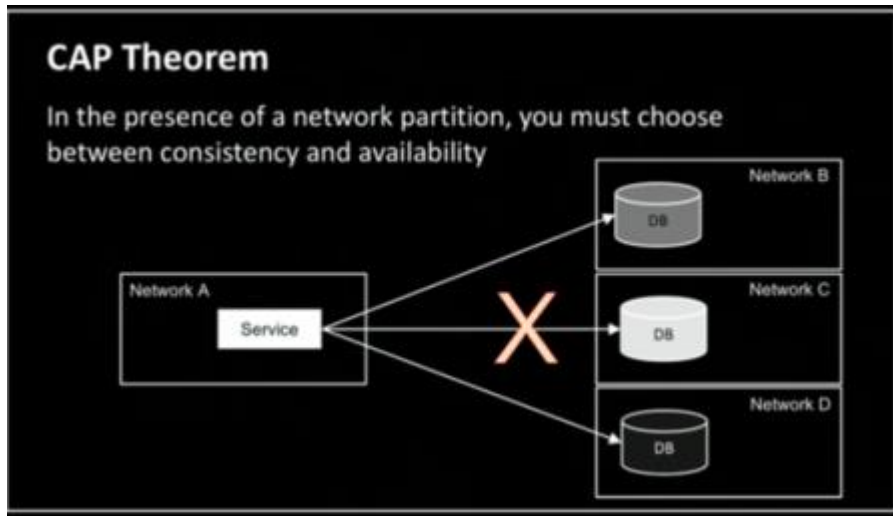
8-9 hours a year



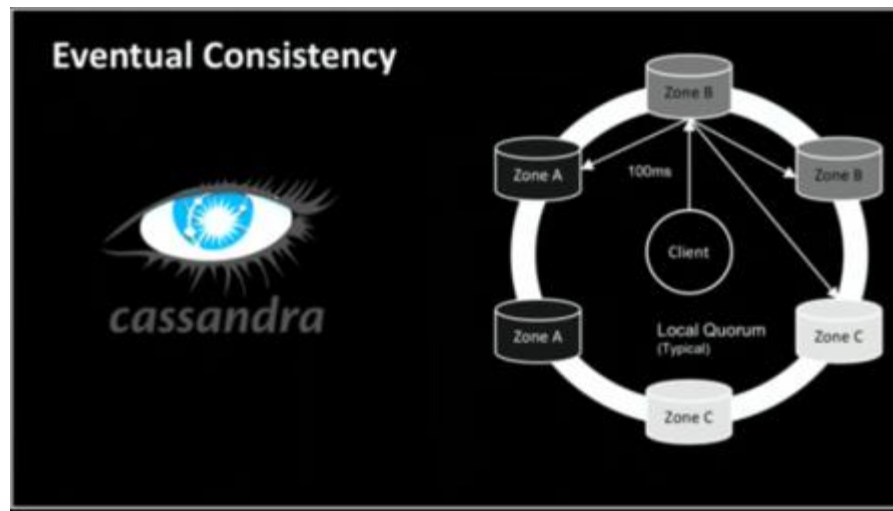
Client libraries are useful



→ simplify client libraries



Service wants to write copy of same data in 3 DBs.  
What if it cannot write in one of them? Fail?



«write to the ones you can get to, and then fix it up afterwards»  
Local quorum

# Dependency: Infrastructure



All in Us East 1



Multi region strategy

- Microservices
- Challenges
  - Dependency
  - **Scale**
  - Variance
  - Change

## What is a stateless service?

Not a cache or a database

Frequently accessed metadata


No instance affinity

Loss a node is a non-event

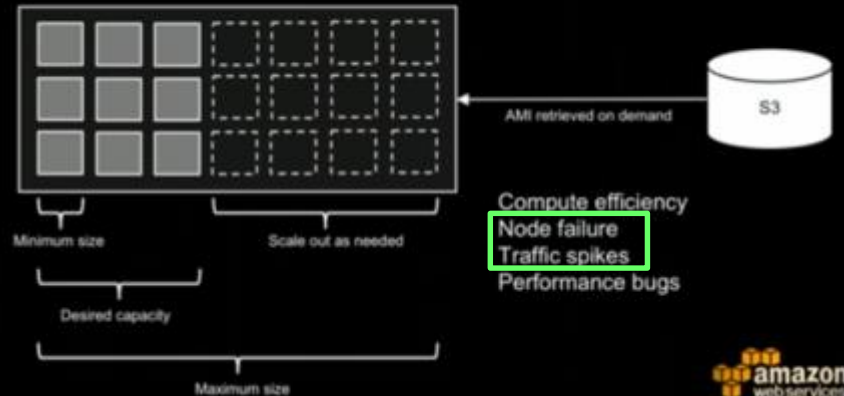
replicate

and test with Chaos Monkey

## Surviving Instance Failure



## Auto Scaling Groups





# Scale: Stateful services

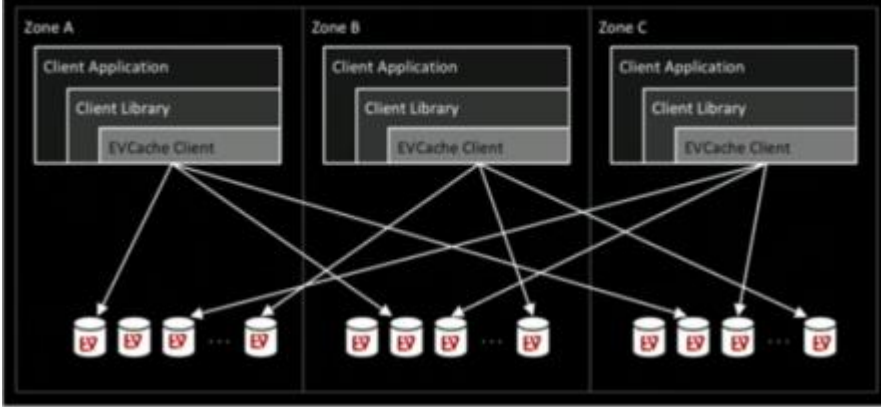
## What is a stateful service?

Databases & caches

Custom apps which hold large amounts of data

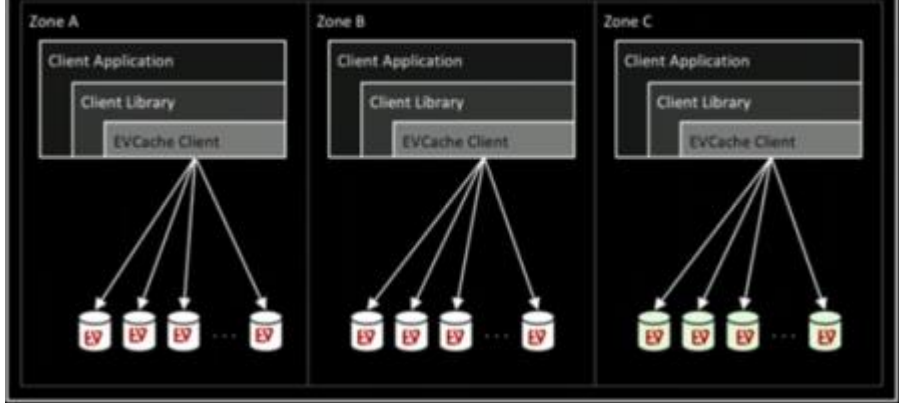
Loss of a node is a notable event

### EVCache Writes



Multiple copies written out to multiple nodes,  
in different availability zones

### EVCache Reads



Local reads, but app can fall back  
to reading across availability zones

30 million requests/sec  
2 trillion requests per day globally

Hundreds of billions of objects  
Tens of thousands of memcached instances

Milliseconds of latency per request    Good scaling!



- Microservices
- Challenges
  - Dependency
  - Scale
  - **Variance**
  - Change

# (Unintentional) Variance: Operational drift

## Operational Drift

### Over time

Alert thresholds  
Timeouts, retries, fallbacks  
Throughput (RPS)

need  
tuning

### Across microservices

Reliability best practices

Best practices (checklist)

## Production Ready

Alerts  
Apache & Tomcat  
Automated canary analysis  
Autoscaling  
Chaos  
Consistent naming  
ELB config  
Healthcheck  
Immutable machine images  
Squeeze testing  
Staged, red/black deployments  
Timeouts, retries, fallbacks


## Continuous Learning & Automation



## (Intentional) Variance: Polyglot & Containers




Engineering clients were going off road



Problems in putting these technologies in the critical path

API gateway integrating scripts that can act as endpoints → monolith pattern

Solution: Push endpoints out of API service  
(into nodejs apps in Docker containers)



## Cost of Variance

Productivity tooling

Insight & triage capabilities

Base image fragmentation

**Node management** → new Titus tier for workload mgmt, autoscaling,  
node replacement ...

Library/platform duplication

Learning curve - production expertise

## Strategic Stance

Raise awareness of costs



Constrain centralized support

Prioritize by impact

Seek reusable solutions



- Microservices
- Challenges
  - Dependency
  - Scale
  - Variance
  - **Change**



# How do we achieve velocity with confidence?



## Production Ready

- Alerts
- Apache & Tomcat
- Automated canary analysis
- Autoscaling
- Chaos
- Consistent naming
- ELB config
- Healthcheck
- Immutable machine images
- Squeeze testing
- Staged, red/black deployments
- Timeouts, retries, fallbacks



## Integrated, Automated Practices



Conformity checks

Red/black pipelines

Automated canaries

Staged deployments

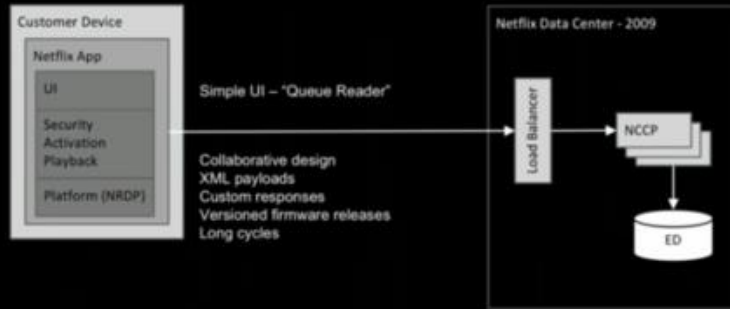
Squeeze tests

live production traffic into new code  
one region at a time



- Microservices
- Challenges
  - Dependency
  - Scale
  - Variance
  - Change
- **Organization & architecture**

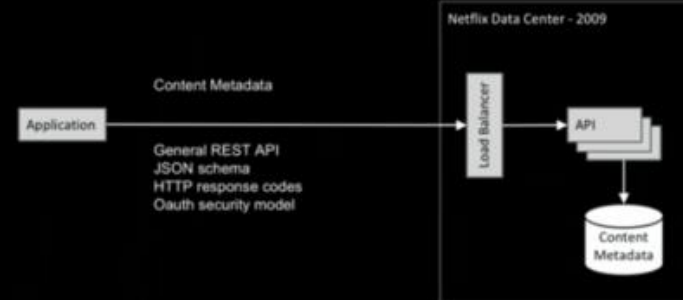
## Electronic Delivery - NRDP 1.x



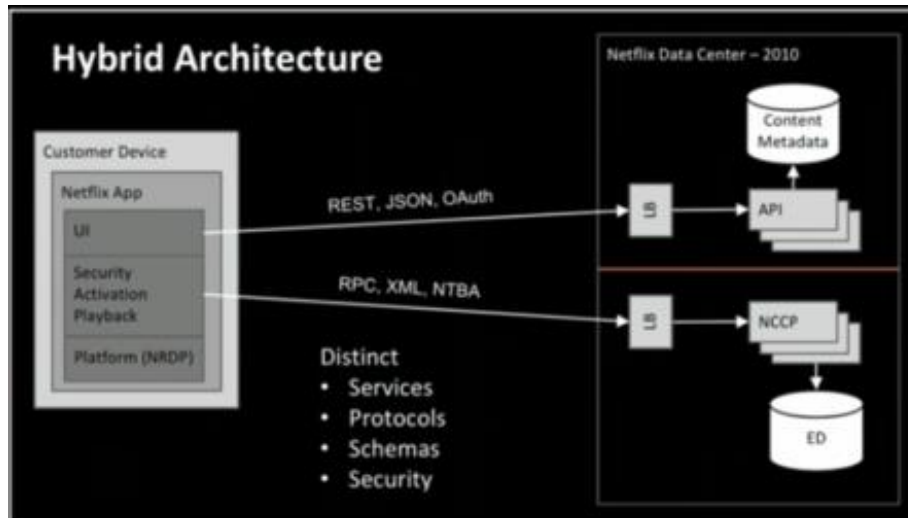
## Netflix API - let a 1000 flowers bloom!



## Netflix API - from public to private



## Hybrid Architecture




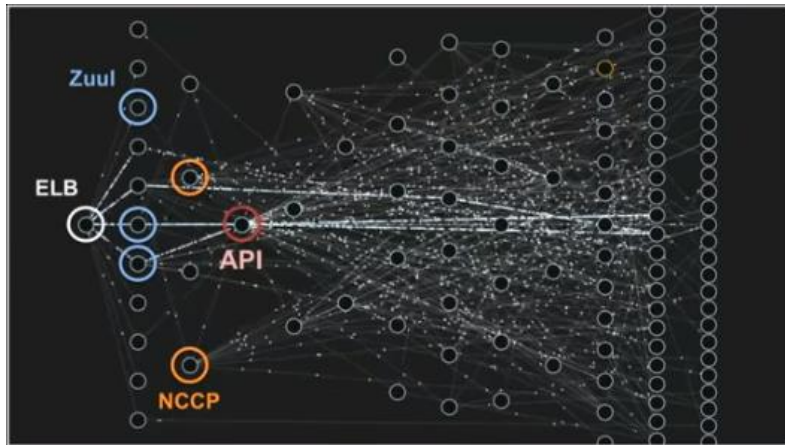
*Josh: what is the right long term architecture?*

*Peter: do you care about the organizational implications?*

## Conway's Law

Organizations which design systems are constrained to produce designs which are copies of the communication structures of these organizations.

Any piece of software reflects the organizational structure that produced it.



NCCP capabilities decomposed into smaller microservices and integrated into Zuul proxy layer and API gateway

## Outcomes & Lessons

### Outcomes

- Productivity & new capabilities
- Refactored organization

### Lessons

- Solutions first, team second
- Reconfigure teams to best support your architecture

# Summary

## Dependency

- Circuit breakers, fallbacks, chaos
- Simple clients
- Eventual consistency
- Multi-region failover

## Variance

- Engineered operations
- Understood cost of variance
- Prioritized support by impact

## Organization & Architecture

- Solutions first, team second

## Scale

- Auto-scaling
- Redundancy – avoid SPoF
- Partitioned workloads
- Failure-driven design
- Chaos under load

## Change

- Automated delivery
- Integrated practices




# NETFLIX

DEVICE PARTNER LIFECYCLE TEAM

- Why microservices?
- Essence of microservices
- Refactoring microservice-based architectures
- Concluding remarks
- Case studies
  - [comparethemarket.com](http://comparethemarket.com)
  - Spotify
  - Netflix
  - Uber

# Uber's initial architecture

- Monolithic architecture built for a single offering in a single city
  - a REST API with which passenger and driver connect
  - three different adapters used to perform billing, payments, sending emails/messages
  - a MySQL database to store all data
- Problems when expanding worldwide
  - all the features had to be re-built, deployed and tested again and again to update a single feature
  - fixing bugs became extremely difficult in a single repository as developers had to change the code again and again
  - scaling the features simultaneously with the introduction of new features worldwide was quite tough to be handled together





# Uber's solution

- Break monolith into microservices
  - introduction of API Gateway through which all the drivers and passengers are connected. From the API Gateway,
  - all the internal points (passenger management, driver management, trip management, etc.) connected from API Gateway
- Individual separate deployable units performing separate functionalities
  - E.g., billing microservices can be deployed independently from the rest
- All features can now scale individually
  - e.g., people searching for cabs > people actually booking & paying cabs

